


INSTITUT  
POLYTECHNIQUE  
DE PARIS



# Checkmate by learning: A modular Reinforcement Learning approach to chess

Enzo Pinchon, Yann Baglin-Bunod, Luis Wiedmann,  
Shashwat Sharma, Mathieu Antonopoulos

# Environment



# Frontend/Backend Architecture

# Backend

- Game Logic Engine: Core chess implementation in backend/src/chess handling rules, moves and state management
- Ai model architecture: modular design where agents can inherit from the base engine class
  - We have diverse AI Implementations using techniques from simple heuristics to advanced algorithms
- Communication Layer: socket-based server to interact with frontend

# Engine & DeepEngine

Inheritance of engines – Modular framework to implement engine and deep engines.

```
with model | generative_head | with_prints | auto_save as env:  
    plot_data = env.train(  
        epochs=epochs,  
        batch_size=batch_size,  
        loader=ld_games | ld_puzzles  
    )  
  
    env.test(loader=ld_games | ld_puzzles)  
  
    env.plot(plot_data)
```

# Frontend Interface

Start a game

New player name

---

Game mode

Player vs AI ▼

AI selection

Random AI - 667 ELO ▼

Player selection

Enzo - 600 ELO ▼

Player 1 color

White ▼

AI selection

Random AI - 667 ELO ▼

Random AI - 667 ELO

AlphaBetaSearchAI - 658 ELO

Greedy AI - 612 ELO

Score CNN2 - 600 ELO

TD Learning AI - 600 ELO

Transformer AI - 600 ELO

Alpha-Beta Transformer - 600 ELO

Tree Search Transformer - 600 ELO

Score CNN - 597 ELO

GreedyExploration AI - 588 ELO

MCTS AI - 588 ELO

Q-Learning AI - 580 ELO

Stockfish AI - 562 ELO

- Frontend allows interactive play and AI benchmarking.
- Supports human vs. AI, AI vs. AI, and human vs. human.
- Visual representation of moves, evaluations, and game states.

# Models


# Model Overview

- **Greedy AI**
- **MCTS AI** (Monte Carlo Tree Search)
- **Random AI**
- **Score CNN**
- **Stockfish AI** (Baseline engine for comparison)
- **Transformer AI**
- **AlphaBetaSearch**



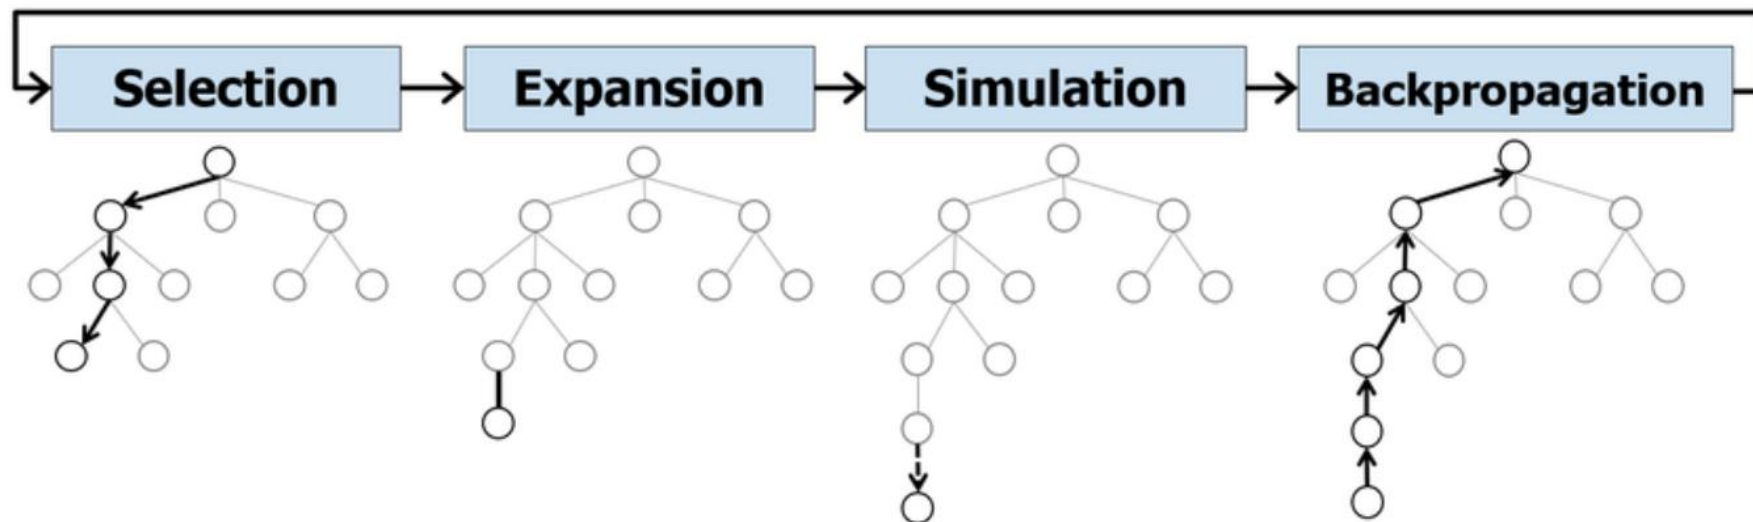
# Greedy / Greedy Exploration

- Optimized Greedy AI that plays as strongly as possible with a single-move evaluation.
- It uses move selection based on piece values (MVV-LVA principle) and positional advantages.
- Strengths:
  - Fast decision-making (no deep search).
  - Simple evaluation based on material and position.
- Limitations:
  - Fails in long-term strategy.
  - Vulnerable to tactical traps.




# MCTS

- Implemented some heuristics and optimizations that guide the search more effectively.
- It uses move selection based on piece values (MVV-LVA principle) and positional advantages .
- Improved evaluation function, weighted random selection explore/exploit.



# Alpha Beta search

- Implemented Alpha-Beta pruning with iterative deepening, inspired by SunFish and Stockfish engines.
- Enhanced pruning efficiency using piece-square tables, move ordering, and time-limited search.
- Evaluation based on material balance, piece mobility, and king safety, without neural networks.



# Deep Networks - Playing

- 1) ChessEmbedding: converts board positions into a high-dimensional latent space.
- 2) GenerativeHead: generates board reconstructions.
- 3) BoardEvaluator: outputs a probability distribution for game outcomes.

**Implemented with both a CNN and a Transformer**

Chess Embedding

One-hot to  
Latent space

GenerativeHead

Ranking list of  
best moves

BoardEvaluator:

White: 78%  
Black: 22%

# Score CNN

- CNN-based evaluation:
  - Captures spatial patterns in board positions.
  - Predicts win probability from a given state.
- Uses convolutional layers with CBAM & SE attention to extract chessboard patterns and enhance feature importance.
  - Incorporates heatmaps to highlight critical board areas for evaluation and move generation.
  - Less effective at long-term planning than transformers but faster and more efficient for local position analysis

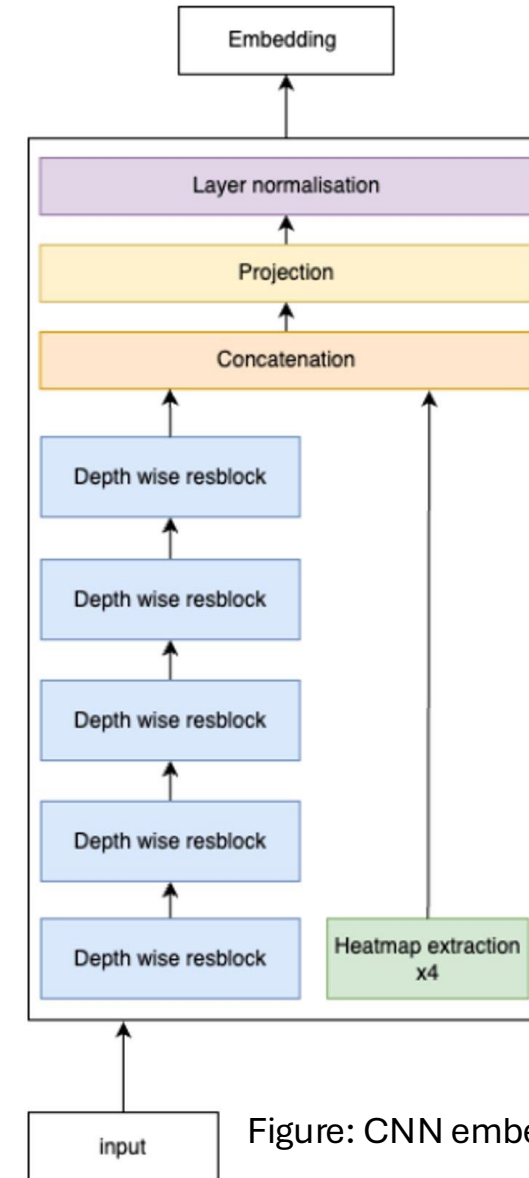
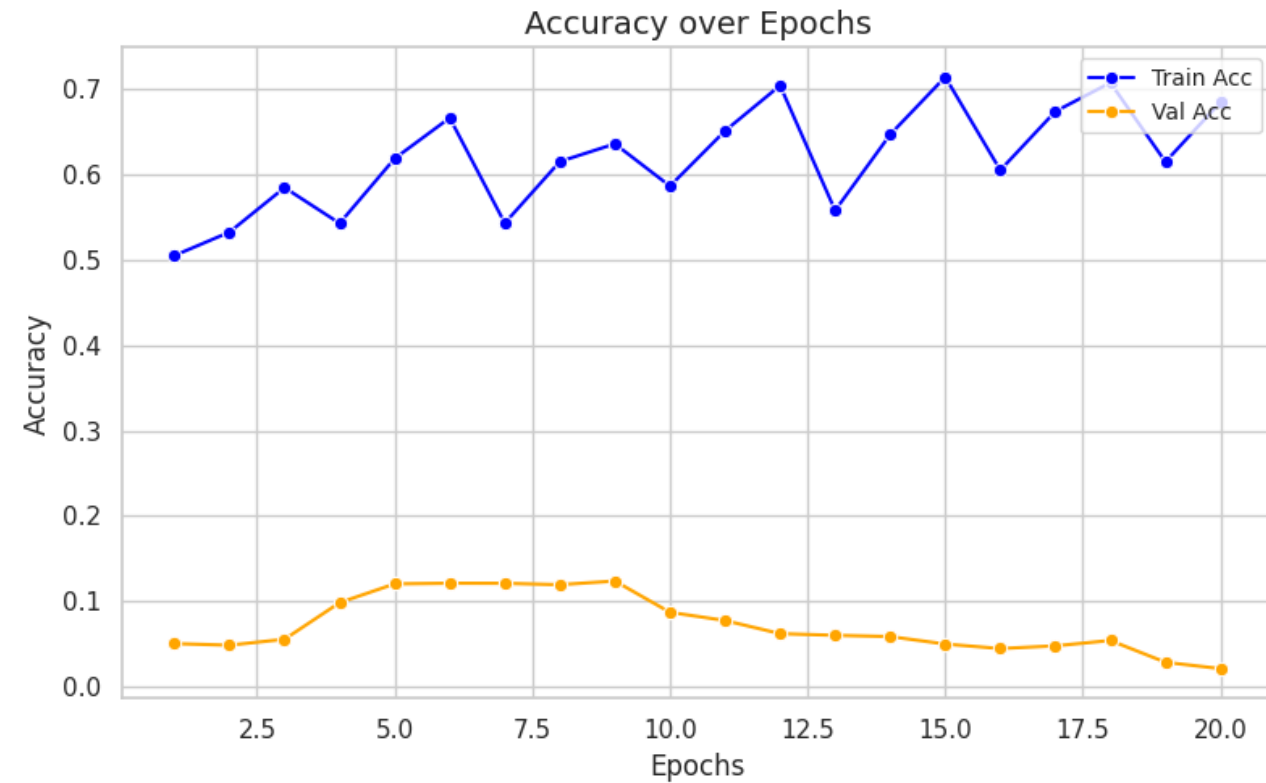
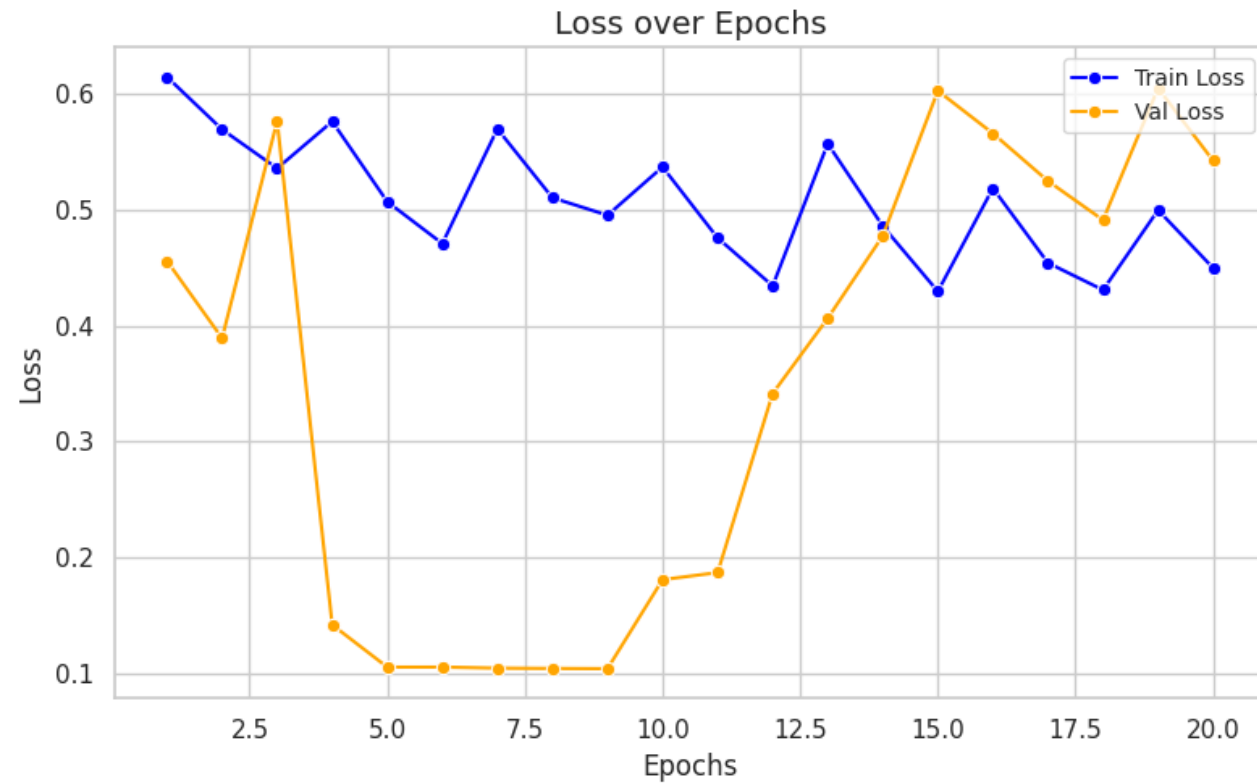


Figure: CNN embedding architecture



# CNN training without attention

Training board evaluation head




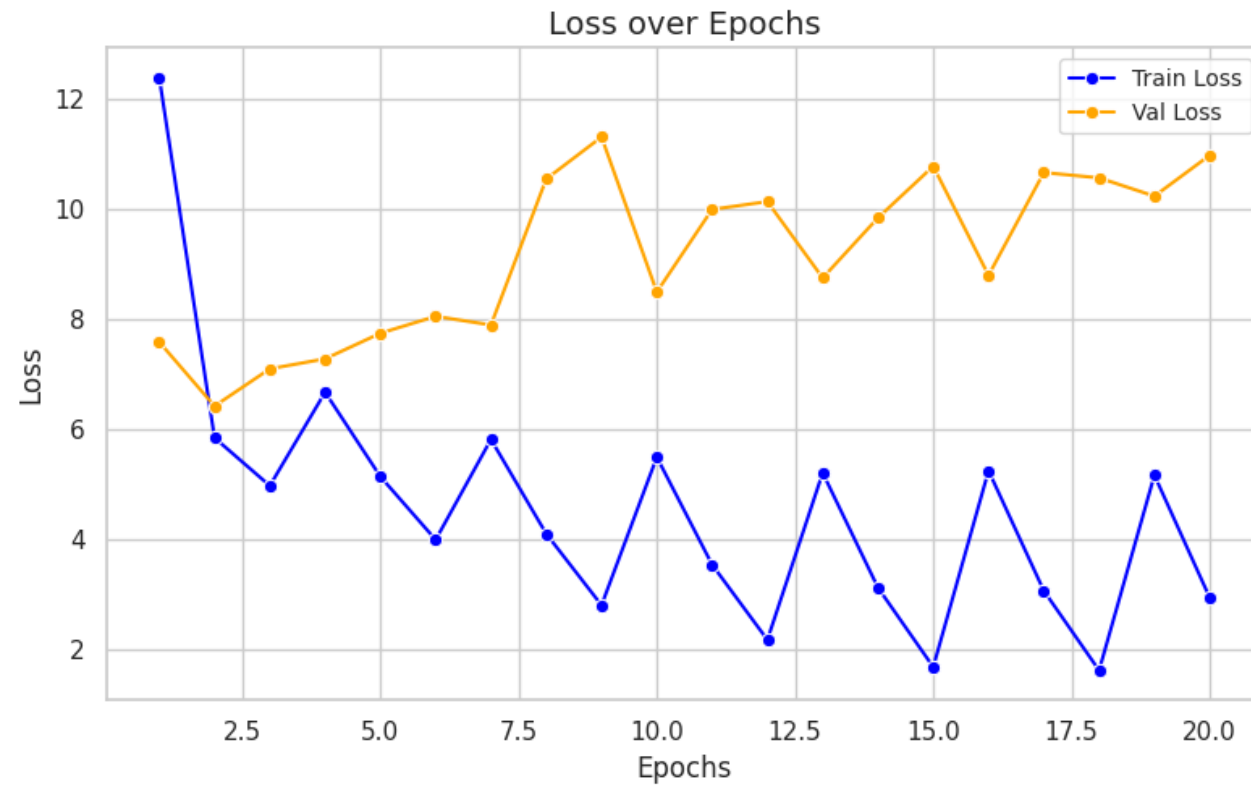
# CNN training without attention

Training generation head



# CNN training with CBAM attention

Training generation head







# CNN training with CBAM attention

Training board evaluation head

Loss over Epochs



Accuracy over Epochs




# Transformer AI

- Transformer-based evaluation:
  - Uses self-attention to evaluate chess positions.
  - Captures long-range dependencies between pieces.
  - Requires large training data & high computation.
- Architecture: Processes  $8 \times 8 \times 13$  board state through a Chess Transformer Encoder (piece + positional embeddings, CNN for local features, transformer blocks for global context).


# Transformer training

Training generation head

Loss over Epochs




Accuracy over Epochs




# Transformer training

Training board evaluation head

Loss over Epochs




Accuracy over Epochs

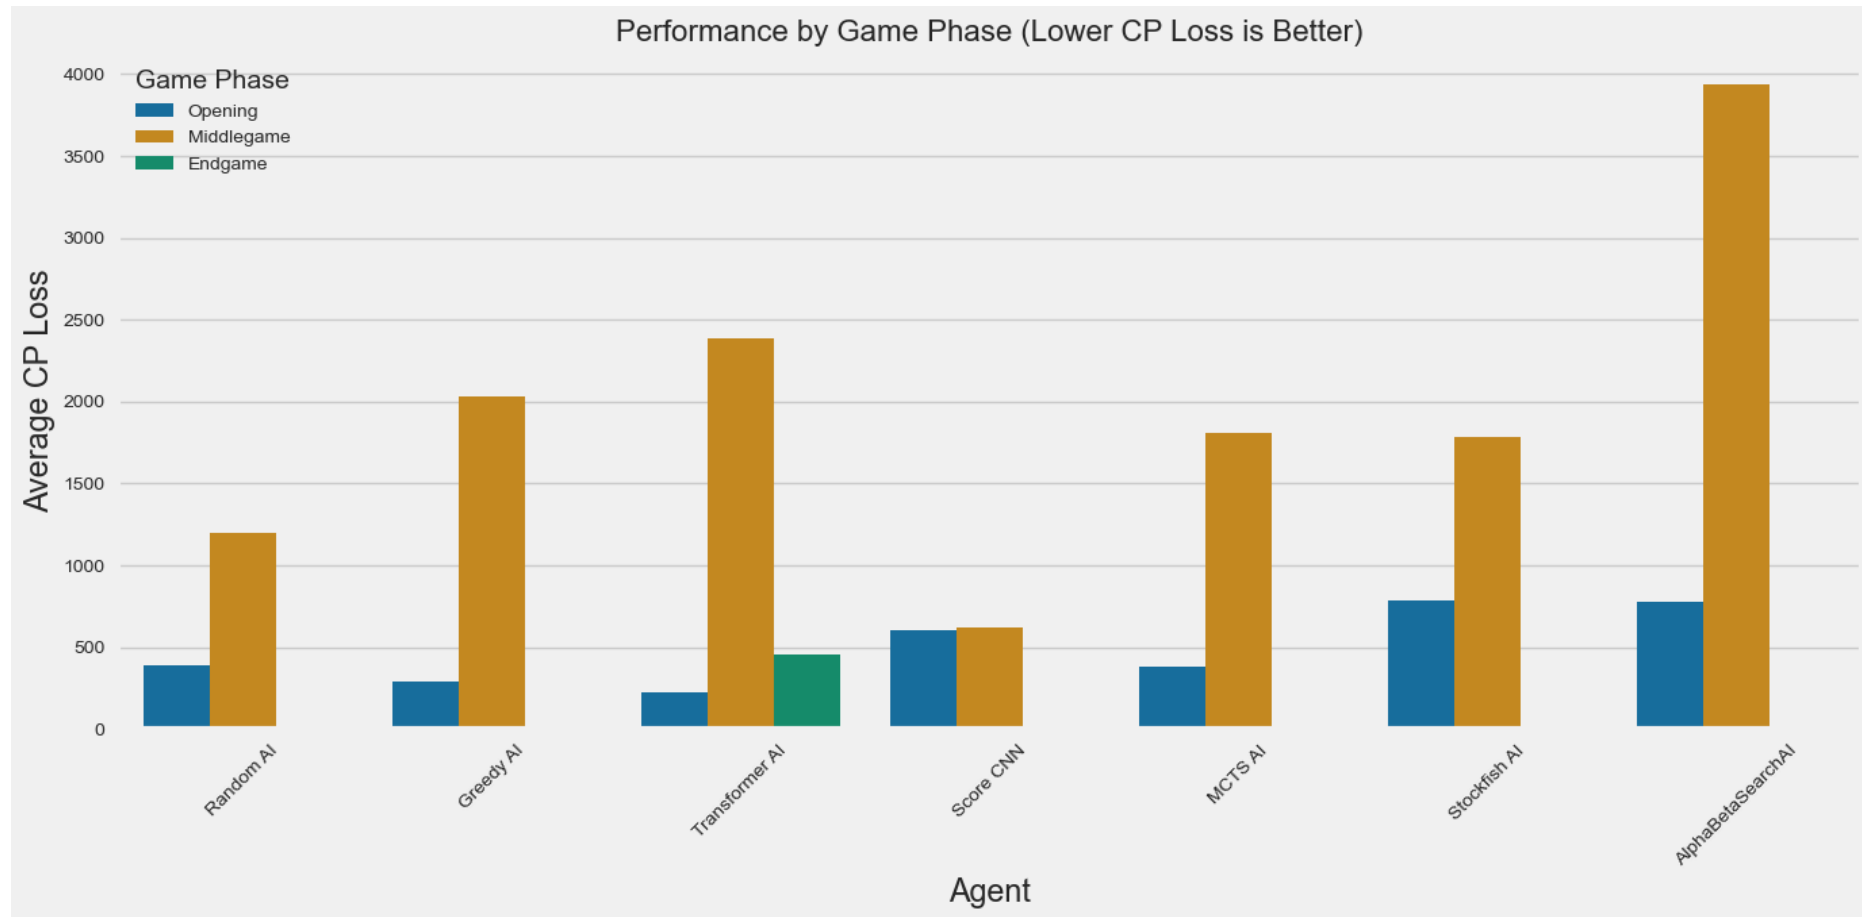


# Evaluation





# Evaluation



# Evaluation





# Evaluation







# Evaluation




# Evaluation



# Evaluation



# Evaluation



# Conclusion

## **Modular Framework**

Enabled rapid experimentation and direct comparison of diverse chess-playing methods.

## **Heuristic Efficiency**

Stockfish and Greedy AI excelled due to efficient search strategies and handcrafted evaluations.

## **Transformer Strengths**

Transformer models effectively leveraged attention mechanisms for global positional understanding.

## **CNN Limitations**

CNN approaches faced generalization challenges related to spatial invariance and sequential dependencies.